

Representing FinReflectKG as a SmartGraph — Options Analysis

Status: v1.0 · 2026-06-30 · for team review

Author: Arthur Keen (ArangoDB)

Scope: how to distribute the FinReflectKG knowledge graph across an ArangoDB cluster as a **SmartGraph** — the options considered, the data evidence behind them, the trade-offs, and the decision.

Sources: measured over the full dataset (DuckDB over the 103 source parquet shards) and verified live on the 3.12.x Enterprise cluster. Backing detail in `sharding-analysis.md` , `multi-distribution-plan.md` §5, `data-analysis.md` , `load-report.md` .

TL;DR

These properties are consequences of how we currently model the source data, not intrinsic facts about it — the graph can be re-modelled (§1, §1.2). Given the current representation:

- **In the source representation, no edge ever connects two companies.** As the data is modelled today (one edge per source triple), every triple is extracted from *one* company's 10-K and stamped with that one `ticker` , so both endpoints of any edge come from the same filing — **0.00% of edges are cross-company, by construction of this representation.** That makes the `filing ticker` the natural smart key and the per-company partition clean.
- **Entities are shared across companies, though** — ~17% of nodes (including ~15% of `PERSON` nodes) are referenced by more than one company. Sharing happens at **shared nodes** that multiple companies independently point to, *not* at cross-company edges (see §1.1). These shared nodes — generic metrics, accounting policies, and generic roles like `board / coo / treasurer` — are the "concept layer," and **27.7% of edges are concept→concept.**
- The shared layer can instead be a **satellite layer** (a hybrid design, Option B): companies form a **disjoint SmartGraph** sharded by `ticker` , **connected to satellite concept collections** replicated to every DB-Server. ArangoDB supports this directly —

a disjoint smart subgraph may connect to satellite collections; it only may not connect to *another* disjoint subgraph (a different `ticker`), which never happens here.

Concept↔concept edges (27.7%) live in the satellite layer. The cost is **replicating the concept layer to every DB-Server** and using **several edge collections** (a relation's two sides must each be all-smart or all-satellite) instead of one relations — which is what surrenders our single-collection fast path, *not* any impossibility.

- **Decision: a Disjoint SmartGraph sharded by `ticker`, with shared concepts duplicated per company** (no satellites). It maximises per-company locality and keeps our single-edge-collection query model (and its fast path) intact. Supernodes are handled by **vertex-centric indexes** (already built on every distribution); Option C adds per-company *locality* on top so no shard holds a hotspot. The cost is bounded node duplication (~2.1×) and aggregating cross-company concept queries by name.

1. What the *current representation* implies (and the choices behind it)

Important caveat — the data does not "force" any of this. We're working with a graph, and there are many valid ways to represent this information. The properties below are **consequences of the modeling choices we inherited from the source**, not intrinsic facts. Specifically, the current graph mirrors the source's row-per-triple form:

1. **one edge per source triple-occurrence** (provenance-preserving; only 8.49 M of the 17.5 M rows are distinct `(entity, relationship, target)` triples — **~52% are repeat observations** across years),
2. **node identity = (name, type) (md5)**, i.e. *no global entity resolution* beyond the source's normalization,
3. **a filing `ticker` stamped on every edge**, and
4. **no cross-filing links** — we never connect two companies' subgraphs beyond what a single filing asserts.

Choices (3) and (4) are *why* the graph looks like a disjoint per-company union and why `ticker` looks like the obvious smart key. **Different choices would change the analysis** — most importantly, **global entity resolution** (one canonical `net income`, one canonical `tim cook`, one canonical company) would create genuine cross-company structure, make the graph *not* disjoint, and push toward global-identity layouts (satellites / Option B) or a different key entirely. Other representations — reified

fact/observation nodes, a separate filer-vs-mentioned-entity layer, role/tenure nodes (§2.2), an n-ary/qualifier model for temporal facts — would each shift the trade-offs. So read the options and decision below as **"given the current representation, here is the best distribution,"** not "the data dictates this." Re-modeling is on the table; see §1.2.

With that framing, the measured shape of the current representation:

Signal	Value	Why it matters for distribution
Edges / nodes / companies	17.51 M / 3.10 M / 743	sizing
Cross-company edges	0.00%	every triple carries one filing ticker , so no edge spans two companies — the per-company partition is clean <i>by construction</i> (not because entities are unique to companies — see §1.1)
Edges owned by the filer ORG	87.2%	the filing ticker is the natural "owner" of an edge
Single-ticker nodes	83%	most nodes belong to exactly one company
Shared (multi-ticker) nodes	17%	the "shared parts" — generic concepts, policies, and roles, referenced by many companies
Edge endpoint classes	company→company 16.0% , company→concept 46.5% , concept→company 9.9% , concept→concept 27.7%	the 27.7% concept→concept decides the shared-layer design (duplicate per company, or put in a satellite layer)
Supernodes	56 nodes with degree >10 K (e.g. net income , in-degree ~99 K)	generic metrics accumulate degree across all 743 companies

Two consequences **(of this representation, not of the data per se)**:

1. **ticker is the natural smart key.** Given (3)+(4) there are no cross-company edges, so sharding by the filing company puts each company's subgraph on one shard and company-scoped queries stay local. Random/hash distribution would scatter every

company across all DBServers. So *for this representation*: a smart key, and the key is `ticker`.

2. **The shared layer is the whole problem.** 17% of nodes and 27.7% of edges involve entities referenced by more than one company. How we place those is the real design question — and it's largely a question of whether we keep the source's per-mention identity or impose global identity (§1.2).

1.1 "But surely the companies share people and entities?" — yes, but the sharing is at *nodes*, not edges

This is the natural objection: multiple years of 10-Ks across 743 companies obviously share directors, officers, regulators, and generic concepts. They do — and we verified it live. The resolution is that **sharing is a property of nodes, not edges**:

- In a sample of 4,000 PERSON nodes, **616 (~15%) are referenced by more than one company**; ~17% of *all* nodes are multi-company.
- The most-shared "people" are **generic roles**, not named individuals (extraction normalizes/lemmatizes names): chief information security officer (162 companies), controller (84), leadership team (82), coo (80), treasurer (66), board (59). These behave just like the shared *concept* layer.

Named individuals show up too — and most are single-company. Spot-checking real executives (live):

PERSON node	# companies	Where
elon musk	1	tsla
michael dell	1	dell
jensen huang	1	nvda
tim cook	2	aapl , nke (Apple CEO + Nike board)


Most named execs are discussed only in their *own* company's 10-K, so they sit entirely within one company's subgraph — no duplication needed. The genuinely cross-company ones are the interesting case: **tim cook appears in both Apple and Nike** (he's on Nike's board). Under Option C he becomes `aapl:tim cook` and `nke:tim cook` — each copy holds that company's edges; "what's Tim Cook connected to at Apple" stays on one shard, and "everywhere Tim Cook appears" aggregates the two copies by name. (Two caveats from the same check: `warren buffett` 's apparent "2 companies" is just the `brk.b` / `brk-b` ticker spellings of *one* company — normalization noise, not real sharing; and `bill gates` has

no node at all — extraction is filing-scoped and Microsoft's 2014–2024 10-Ks don't name him, so former founders can be absent.)



Worked example — Tim Cook, today vs. Option C (the chosen disjoint design)

His actual records (live): one `tim cook` `PERSON` node with 5 inbound edges — 4 from Apple (`depends_on` 2022/2023/2024, `subject_to` 2016) and 1 from Nike (`appoints` 2018). What ties each edge to a company is its `_from` (the company `ORG` node) and its `ticker` field — there is no Apple→Nike edge; the two companies only meet *at* this node.

Today — FinReflectKG : one shared node, both companies point at it



Option C — Disjoint SmartGraph (sharded by `ticker`): one copy per company, each local to its shard



The single node splits into `aapl:tim cook` and `nke:tim cook` (`key = <ticker>:<md5(name|type)>`); Apple's 4 edges route to the Apple copy and Nike's 1 edge to the Nike copy, each fully on its company's shard. Nothing is lost — each edge already carried exactly one `ticker`, so it maps cleanly to one copy. "Tim Cook at Apple" is a one-shard lookup; "everywhere Tim Cook appears" resolves both copies **by name** and unions them.

The general rule (beyond Tim Cook): a shared node like `board` is pointed at by edges from `jpm`, `gs`, `bac`, ... but **each carries its own single ticker** — there is no `jpm → gs` edge. **Companies only ever meet at a shared node, never across an edge.**

That is the property the options in §2 exploit (by duplicating or replicating those shared nodes) — and, per the caveat above, it holds only for the current representation (§1.2).

1.2 We are not locked into the source's triple representation

The whole analysis above assumes the current row-per-triple model. We can re-model the data; doing so changes the distribution question. The main alternatives:

Alternative representation	What changes	Effect on distribution
Global entity resolution (canonicalize net income , tim cook , companies to one node each)	shared entities become genuinely shared, bridging companies	graph is no longer disjoint → ticker duplication (Option C) loses canonical identity; pushes toward global-identity layouts (satellites / Option B) or a non- ticker key
Reified fact / observation nodes (a Fact node per distinct (e,r,t) , with per-filing- year observation edges)	dedups to 8.49 M distinct facts (~52% of rows are repeats); time/provenance move onto observations	new node/edge counts; a natural home for temporal data; changes what "an edge" means for sharding
Filer-vs-mentioned- entity split (separate the 743 filers from the entities they mention)	clarifies ownership; could make filers smart and mentions satellite	a principled basis for Option B's company/shared split
Role / tenure nodes (§2.2)	reify person↔company appointments with validFrom / validTo	lets people stay single global (satellite) nodes; targeted temporal model
n-ary / qualifier model (facts with date/source qualifiers as first-class)	richer temporal & provenance semantics	heavier; only if temporal/provenance querying becomes central
Harmonizing <i>additional</i> sources — other structured datasets, and/or GraphRAG	new data may not be filing/ ticker -scoped, may introduce entities with no single	erodes the "every edge has one ticker " and "disjoint" assumptions → can force global-

Alternative representation	What changes	Effect on distribution
extraction over new unstructured text	company owner, and may add genuine cross-company links	identity layouts or a different/compound key

This dataset is unlikely to stay a single, homogeneous source. It is plausible that other data will be **harmonized into** it over time — additional **structured** datasets, and/or facts extracted from **unstructured text via GraphRAG import**. That matters here because the entire "disjoint by `ticker` " argument rests on a property of *this* source: every triple is filing-scoped and carries one company `ticker` . Newly harmonized data may carry **no `ticker`** , may describe entities **shared across companies**, or may assert **direct company-to-company links** — any of which **breaks disjointness** and weakens `ticker` as the smart key. So the distribution choice should be treated as **revisable as sources are added**, and a design that depends hard on strict disjointness (Option C) is more brittle to source harmonization than a **global-identity layout** (entity resolution + satellites/Option B), which absorbs heterogeneous, cross-cutting data more gracefully.

Why Option C is still the right call now: the immediate requirement is a SmartGraph with **text co-located with the entities it relates to**, which is inherently per-company — so keeping the source's per-mention identity and sharding by `ticker` fits *that* use case and is the cheapest path (no entity-resolution project). The honest framing: **Option C is the best choice given the current representation and the current use case — not because the data dictates it.** If/when additional sources are harmonized in, or the priority shifts to global cross-company analytics or a canonical knowledge graph, entity resolution + a global-identity layout becomes the better foundation, and we revisit.

2. The options (given the current representation)

Note on the smart key: the smart key in every SmartGraph option below is `ticker` — i.e. **the company's stock symbol**. With 0% cross-company edges and a `ticker` on every edge, the stock symbol is the natural and chosen smart key; the options differ only in how they handle nodes that have *no single* stock symbol (the shared concepts and people).

Option A — Random / hash distribution (no smart key) ❌

Shard everything by `_key` . Simple, no modelling. But it ignores the disjoint structure: a company's nodes and edges scatter across all shards, so the typical "expand this company" query hits every DBServer. **Rejected** — it throws away the one property (no cross-company edges) that makes this graph cluster-friendly.

Option A' — Non-disjoint SmartGraph by stock symbol, *no* duplication ❌

Use `ticker` (stock symbol) as the smart key but **don't** duplicate shared nodes — assign each shared concept a single (arbitrary) stock symbol and leave it there. The catch: edges from *other* companies to that concept then have **mismatched smart values on their two endpoints**, so they become **cross-shard** in a non-disjoint SmartGraph. Since **company→concept is 46.5% of all edges**, this scatters nearly half the graph — losing locality exactly on the heavily-referenced shared layer. **Rejected** — "smart key = stock symbol" only delivers locality when paired with a shared-node mechanism (duplication → Option C, or satellites → Option B); on its own it doesn't.

Option B — Disjoint SmartGraph (companies) connected to a satellite concept layer ⚠️

The "SmartGraph with satellites for the shared parts," done the way ArangoDB actually supports it — a **disjoint company SmartGraph hanging off a satellite concept layer**:

- Node = **smart** collection, sharded by `ticker` , `isDisjoint: true` (the 83% company-owned nodes; each company is its own disjoint subgraph).
- SharedNode = **satellite** collection (the 17% shared concepts — one global identity, replicated to every DB-Server, read-local everywhere).
- **company↔concept edges (smart↔satellite, 56.4%)** — fully supported; "all smart nodes can be connected to satellite nodes." Sharded by the company; the satellite endpoint is local everywhere.
- **concept↔concept edges (satellite↔satellite, 27.7%)** — supported; they live entirely **within the satellite layer** (a satellite edge collection, replicated everywhere).
- **company↔company edges (16%)** stay within one `ticker` — fine. The *only* prohibition is connecting **two different** disjoint subgraphs (different `ticker` s) directly, and there are **0%** of those.

The one structural consequence: because a relation's two sides must each be type-homogeneous (all-smart or all-satellite), the edges can't all sit in one

relations collection — they split into **several edge collections by endpoint type** (company→company, company→concept, concept→company, concept→concept). (The earlier spike that errored — "...required to be a smart collection. But would be created as satellites." — was an attempt to keep one mixed-type edge collection; the supported form is the layered split above.)

Net: fully feasible and a clean layered design. Pros: **no duplication**; concepts keep **one global identity**; concept reads are local on every DB-Server; cross-company analytical traversals (§3.1) keep working. Cons: the concept layer — its **nodes (17%) and the concept→concept edges (27.7%)** — is **replicated to every DB-Server** (verified live; company↔concept edges are smart-sharded, not replicated), and the edge model becomes **several collections** instead of one relations , which costs us the single-collection VCI direct-edge fast path (§3 finding (a)).

Option C — Disjoint SmartGraph, concepts duplicated per company (no satellites) chosen

- Smart key = ticker ; node key = <ticker><hash(name,type)> .
- A concept shared by N companies is **duplicated into each of those N companies' shards**, so every node has an owning company and **every edge is intra-shard** — including the former concept→concept edges. Fully **disjoint** (isDisjoint: true).
- One Node collection, one relations collection — unchanged from today.
- **Cost:** nodes grow to ~6.66 M (**2.1x**; edges are *not* duplicated). A concept like net income becomes ~743 small nodes instead of one, so "every company that discloses net income" must aggregate the copies **by name** (cheap — there's an index on name).

2.1 Collection mapping — Option B vs Option C (these are *not* the same collections)

The two designs have **different collection sets**, and where a name is reused (Node , relations) it holds **different membership**. This is the part most likely to be conflated, so it is spelled out explicitly.

Option C — Disjoint SmartGraph: 1 vertex + 1 edge collection (+ chunks).

Collection	Kind	Sharding	Contents
Node	smart vertex (smartGraphAttribute=ticker)	by ticker	all nodes — company-owned <i>and</i> shared, shared ones

Collection	Kind	Sharding	Contents
			duplicated per referencing ticker (~6.66 M docs)
relations	smart edge	by ticker	all 17.5 M edges, incl. concept→concept
chunks	smart orphan (or distributeShardsLike: Node)	by ticker	source text, co-located

Option B — Disjoint SmartGraph + satellite layer: 2 vertex + several edge collections (+ chunks).

The edges split by endpoint type because a relation's two sides must each be all-smart or all-satellite:

Collection	Kind	Sharding	Contents
Node	smart vertex	by ticker	only company-owned nodes (83%), single copy each
SharedNode	satellite vertex	replicated to every DBServer	only shared nodes (17%), one identity, physically on all DBServers
rel_company	smart edge	by ticker (distributeShardsLike the smart Node)	company→company (smart→smart, 16%)
rel_company_concept / rel_concept_company	smart edge	by ticker (distributeShardsLike the smart Node)	company↔concept (smart↔satellite, 56.4%); the satellite concept endpoint is local everywhere, so the hop is still local
rel_concept	satellite edge	replicated to every DB- Server	concept→concept (satellite→satellite, 27.7%)
chunks	satellite or smart	—	source text

Verified live (throwaway hybrid disjoint SmartGraph on the 3.12.x cluster): the
smart↔satellite edge collection comes up with distributeShardsLike the smart vertex

collection and replicationFactor: 1 — i.e. **smart-sharded, *not* replicated**. Only the satellite↔satellite edge collection comes up replicationFactor: "satellite" . So under Option B the replication cost is the **concept vertices (17% of nodes) + the concept→concept edges (27.7%)** — *not* the company↔concept edges.

Logical entity → collection:

Logical entity	Option C	Option B
Company-owned node (1 ticker, 83%)	Node (smart)	Node (smart)
Shared node (≥2 tickers, 17%) — concepts, board , Tim Cook	Node (smart), duplicated → N docs	SharedNode (satellite), 1 doc replicated
company↔company edge (16%)	relations (smart)	rel_company (smart)
company↔concept edge (56.4%)	relations (smart)	rel_company_concept / rel_concept_company (smart↔satellite)
concept↔concept edge (27.7%)	relations (smart) — endpoints duplicated, so ticker-local	rel_concept (satellite edge, replicated)
Source text	chunks (smart by ticker)	chunks

Where they explicitly do *not* correspond:


1. Option C has **one** edge collection (relations) for *all* edges and **one** vertex collection (Node). Option B has a separate satellite vertex collection (SharedNode) and **splits edges into several collections by endpoint type** — none of which exist in Option C.
2. Node membership differs: Option C Node = *everything* (company + duplicated shared); Option B Node = *company-owned only* (shared nodes live in SharedNode).
3. Option C's single relations holds all edges; Option B has no single edge collection at all — edges are partitioned by (from-type, to-type).
4. Option B needs two classification steps Option C does not: label each node company-owned vs shared (→ Node vs SharedNode), and route each edge by its endpoints' types (→ the right edge collection).

Tim Cook, concretely:

- **Option C:** two docs in **Node** — Node/aapl:56d4... , Node/nke:56d4... ; his 5 edges in **relations** .
- **Option B:** one doc in **SharedNode** — SharedNode/56d4... (no prefix), replicated to every DBServer. His 5 edges are company→person = **smart↔satellite**, so they live in the **re_l_company_concept** (smart→satellite) edge collection — **not** the concept→concept satellite collection (re_l_concept is strictly satellite↔satellite).

2.2 Refinement (for the satellite option): reified role / tenure nodes

A way to make the **satellite option (B)** elegant for *people* — keeping an individual as a **single** global node instead of duplicating them — is to **reify the appointment as a node**. Instead of company → person edges, introduce a company-owned **role / tenure node** carrying validFrom / validTo , that references the (satellite) person:



Why it works with the distribution: the role node is **company-owned (smart, sharded by stock symbol)** and the person is the **satellite**, so the role → person edge has **one smart endpoint** — so it lives in the smart edge collection (sharded by the role's stock symbol; the person's replica is local). This means people **never need the separate satellite edge collection** that concept→concept edges force, because there's never a satellite-to-satellite edge.

The data supports it (measured live): POSITION -typed nodes already exist (**9,984**); PERSON nodes already use role verbs — serves_as (21K), holds_position (19K), served_as (13K), member_of (9K), works_for (6K), held_position (3.6K); and **~99% of PERSON edges already carry a startDate** , so the from-dates are present.

What it buys: single global identity for the individual with **no duplication**;
clean **role-tenure** modeling (`validFrom` / `validTo`) — a *targeted* temporal model
exactly where identity-over-time matters, far cheaper than the full time-travel
blueprint (§6).

Caveats:

- **Only solves the PERSON/role dimension.** The 27.7% concept→concept edges (metric↔metric, policy↔policy, ...) are not role-based, so they remain satellite→satellite — Option B's edge split is narrowed, not eliminated.
- **It's an ETL enrichment, not free:** the source has POSITION nodes, role verbs, and dates, but not clean reified `person-role-company-dates` appointments. The role vocabulary is messy (past/present variants; titles typed inconsistently as PERSON / ORG / POSITION), so synthesizing role nodes carries normalization risk.
- **Changes the schema/query shape** to a 2-hop `person → role → company` pattern.
- **Not needed by Option C (chosen):** there, Tim Cook is simply duplicated and the dates already live on the edges. This refinement matters if we revisit the satellite option, or if "track an individual's roles across companies over time" becomes a first-class requirement.

3. Key technical findings

Three findings drive the recommendation: (a) the edge model / fast path, (b) supernode handling, and (c) the query workload (§3.1).

(a) Option B fragments our fast path; Option C preserves it.

On this 3.12.x cluster we verified that the fast access path is the **direct edge-collection query**

(`FOR e IN relations FILTER e._from == ... AND e.type == ... AND e._toType == ...`), which uses our vertex-centric indexes and narrows perfectly — *not*

pattern traversals, which fall back to the generic edge index. Our entire query layer (benchmark queries, visualizer actions, Cypher→AQL) is built on a single `relations` collection. Option B's split would force those direct queries to `UNION` several edge collections, each with its own indexes — a permanent tax on the path we proved is fast. Option C leaves `relations` whole.

(b) Vertex-centric indexes are the primary supernode tool; Option C adds locality on top. To be clear: the supernodes are handled first and foremost by **vertex-centric**

indexes (VCIs), which we already build on every distribution — `relations(_from, type, _toType)` and `relations(_to, type, _fromType)` . On the `net income` supernode (in-degree ~99 K), the typed **direct edge query** ("which ORGs disclose net income") scans only the **59,315 matching edges** via the VCI instead of all ~99 K. That selectivity win is independent of the distribution choice — it applies to the baseline, OneShard, and either SmartGraph option. (Caveat, per `load-report.md` : the VCI engages for *direct edge queries*, not pattern traversals, so supernode-touching queries must be written as direct edge queries.)

What the **distribution** choice changes is a *different* axis — **locality** (how the supernode's edges are spread across shards), not selectivity:

- Under **Option C**, `net income` is duplicated per company, so each copy carries only one company's edges. A company-scoped supernode query stays on a single shard, and no shard holds a 99 K-edge hotspot. (A *global* "all companies disclosing net income" still aggregates across shards and by name.)
- Under **Option B**, `net income` stays one node whose inbound edges live in the company-sharded smart collection — so a global scan still scatters across the cluster (the VCI helps *within* each shard, but you touch all shards).

So VCI (selectivity) and Option C's per-company decomposition (locality) **stack**; the decomposition is a complement to the VCI, not a replacement for it.

3.1 Query workload vs. distribution — checking the 22 example queries (`cypher-queries.md`)

Distribution can't be judged in the abstract, so we checked Option C against the 22 representative queries in `cypher-queries.md` . The decisive mechanic: **under Option C a node is duplicated per ticker , so a multi-hop path survives only if every edge on it shares the same ticker** . A path that bridges two companies through a shared node (e.g. company B appears in CINF's filing *and* in B's own filing) lands on `cinf:B` for one hop and `B:B` for the next — **different documents — so the traversal breaks** (or needs a name-based re-resolution + shard hop at each company boundary). A live probe confirmed cross- `ticker` 2-hop paths are common (risk→dependency was 3/3 cross-ticker in the sample).

Bucket	Queries	Behaviour under Option C (disjoint by <code>ticker</code>)
Global aggregation	1, 2 (type distributions)	distribution-agnostic; scatter-gather on any layout

Bucket	Queries	Behaviour under Option C (disjoint by ticker)
Company-scoped / within-filing	3, 4, 5, 13, 15	clean & local — single-company root, all edges share that ticker
Cross-company multi-hop bridges	6, 7, 8, 9, 10, 11, 12, 14, 16, 17, 18, 19, 20, 21, 22	survive only where the path stays within one ticker ; otherwise break and need rewriting (name re-resolution at each company hop). Dependency cycles (18, 22) and stakeholder→company→metric (6, 7, 9) are inherently cross-filing.

So ~15 of the 22 — the richer analytical / GraphRAG traversals — depend on **global node identity** to stitch facts across companies' filings, which is exactly what Option C's duplication removes. This is the **opposite** pressure from finding (a): Option B's single global identity *preserves* those bridges, while Option C *breaks* them.

Resolution — this is a *why we have multiple distributions* point, not an Option C defect:

- The **cross-company analytical suite** should run on **OneShard (or the baseline)**, where global identity is intact and (on OneShard) there's no cross-DBServer penalty.
- The **disjoint SmartGraph** serves the **company-scoped, text-co-located** workload — where queries are rooted at one company and stay local.
- If that analytical suite is ever needed *on* the SmartGraph, those ~15 queries must be rewritten with name-based re-resolution at company boundaries — a real, documented cost.

4. Comparison

Dimension	A: Random	B: Hybrid + Satellites	C: Disjoint + duplication 
Per-company locality	poor	good	best
Edge model	single	several (split by endpoint type)	single (unchanged)
Preserves VCI direct-edge fast path	n/a	no (union needed)	yes

Dimension	A: Random	B: Hybrid + Satellites	C: Disjoint + duplication 
Supernode <i>selectivity</i> (VCI — all builds)	yes	yes	yes
Supernode <i>locality</i> (distribution)	hot	hot (global scatter)	decomposed per company
Concept identity	global	global	split (aggregate by name)
Company-scoped queries (1-hop, within-filing)	scatter	local	local
Cross-company multi-hop traversals (≈ 15 of the 22, §3.1)	work (global identity)	work (global identity)	break / need name re-resolution → run on OneShard
Node storage	3.1 M	3.1 M (17% replicated)	~ 6.66 M ($2.1\times$)
Edge storage	17.5 M	17.5 M (concept→concept 27.7% replicated $\times N$; rest smart-sharded)	17.5 M
Build complexity / risk	low	high (2 vtx + several edge collections, hybrid rules)	medium
Is it a SmartGraph?	no	yes	yes

5. Decision & rationale

Chosen: Option C — Disjoint SmartGraph sharded by ticker , shared concepts duplicated per company.

The driving requirement (from the downstream text-to-graph consumer): a SmartGraph is a **hard requirement of the text-to-graph tooling**, and the goal is to **co-locate the text-to-graph elements with the entities they relate to** — there is no requirement for SatelliteCollections specifically, and no current need for global single-identity

concept analytics. Given that, Option C is both the right fit and the better engineering choice:

- It **is** a SmartGraph (satisfies the tooling requirement).
- **Best co-location**, including the source text: we also smart-shard the `chunks` (source-text) collection by `ticker`, so a company's nodes, edges, *and* the text they were extracted from all land on the same shard. (A satellite would *replicate* 1.4 M chunks to every DBServer instead of co-locating them — the opposite of the requirement.)
- Keeps the single- `relations` model → the VCI fast path and **company-scoped** queries work unchanged. (Cross-company *multi-hop* queries are the exception — see §3.1; those belong on the OneShard/baseline graph.)
- Adds supernode *locality* on top of the VCIs: the per-company copies mean no shard holds a hotspot (VCIs already give the selectivity on every build).
- Disjointness is real (0% cross-company), so duplication is the only cost — bounded at 2.1× nodes, and edges don't duplicate (nodes are ~100 B, so it's cheap).

Accepted trade-off: concept identity is split across companies, so global cross-company concept queries (“which companies disclose metric X?”) aggregate BY `name`. This is cheap with the `name` index, and we'll monitor it.

Monitoring / when we'd revisit: after the build we check the duplication ratio (~2.1×, edges not duplicated), shard balance across the 743 tickers (watch the big filers like `etr` / `pru` / `met`), and query locality (`explain` shows no cross-DBServer `RemoteNode` on company-scoped queries). The only trigger to reconsider Option B is if global cross-company concept rollups turn out to dominate the workload — in which case we'd either switch to the satellite layer or materialise a concept-aggregate view.

6. Related question — the data spans 2014–2024, so do we need the time-travel blueprint?

A fair question that comes up because the 10-Ks span 11 filing years: should we model this with ArangoDB's **time-traveling graphs blueprint** (immutable proxy + versioned documents + `created` / `expired` intervals + `hasVersion`, traversed "as of T")?

Recommendation: no — keep temporal data as edge properties. The blueprint solves a different problem than the one this data presents.

What the data is: an append-only observation log, not versioned entities

The blueprint is designed for **mutable entities that evolve with supersession** (e.g. a CMDB device whose config changes — you reconstruct the single valid version as of an instant). FinReflectKG is the opposite shape: a pile of **dated observations**, each extracted from one annual filing and stamped with its `year` + extracted `startDate` / `endDate` . The same fact is **re-asserted every year, not replaced**. Live example — `aapl --discloses--> net income , edges` per filing year:

	2014	2015	2016	2017	2018	2019	2020	2021	2022	2023	2024
edges	18	17	8	16	9	11	9	10	9	7	6

It appears in **every** filing year. There is no "version N supersedes N-1" relationship between these — they're independent, overlapping assertions (this is also why **~52% of all rows are repeat observations** across years — only 8.49 M of 17.5 M rows are distinct triples). Proxy/version modelling would impose a supersession semantics the data doesn't have.

The temporal dimension is already first-class — on the edges

Every edge carries `year` , `startDate` / `endDate` (ISO YYYY-MM , sortable), `startDateRaw` / `endDateRaw` , with a `rel_ticker_year` index. ~95% of edges have an explicit `endDate` validity window (~4.5% are open-ended / `default_end_timestamp`). The natural temporal questions are plain filters:

- "what did company X report in filing year Y" → `FILTER e.year == Y`
- "relationships valid during a window" → `FILTER e.startDate <= @t AND (e.endDate == null OR e.endDate >= @t)`
- "how a relationship trended year over year" → `COLLECT BY year` (the table above)

Query 19 in `cypher-queries.md` already does time-ordered impact chains this way.

Cost / fit

The blueprint would explode 17.5 M edges into proxies + version documents + `hasVersion` edges, require `WITH` clauses and as-of traversal logic throughout, and **work against the disjoint-SmartGraph plan** (more collections to co-locate per company) — a large remodel for a poor semantic match.

When we'd revisit

Only if a genuine **versioned-snapshot** requirement appears — "reconstruct the full entity-state graph as of an *arbitrary* instant, with supersession," at a scale where property filters don't suffice. The proportionate upgrade short of the blueprint is a **temporal index** (`relations(year, startDate)`) + a documented set of "as-of" AQL patterns. The current per-filing grain already supports as-of-**year** reconstruction, which matches how the data arrives (one annual filing at a time).

Appendix — headline numbers

- 17,513,372 edges · 3,099,773 nodes · 1,384,513 chunks · 743 companies.
- **0.00% cross-company edges** (every triple carries one filing ticker); 87.2% of edges owned by the filer ORG.
- Nodes: 83% single-ticker, **17% shared across companies** (entities *do* repeat — sharing is at nodes, not edges). PERSON sample: 616/4,000 (~15%) multi-company; most-shared are generic roles (`board` , `coo` , `treasurer` , ...).
- Duplicated model $\approx$ 6,658,668 nodes (2.1 $\times$).
- Edge endpoint mix: company→company 16.0%, company→concept 46.5%, concept→company 9.9%, concept→concept 27.7%.
- Supernode example: `net income (FIN_METRIC)` in-degree ~99 K; typed direct query scans 59,315 matching edges via the (`_to` , `type` , `_fromType`) VCI (selectivity, on every build). Option C additionally decomposes it into ~743 per-company copies (locality).

Verification status

Every key figure and ArangoDB-behavior claim in this document, and how it was checked (2026-06-30):

Re-derived from the source parquet (DuckDB, this review):

- 17,513,372 edges · 3,099,773 nodes · 743 tickers.
- 83% single-ticker / **17% shared** nodes (2,572,439 / 527,334).
- Duplication factor **2.15 $\times$** (6,658,668 node copies).

- Endpoint mix **16.0 / 46.5 / 9.9 / 27.7%** (company→company / company→concept / concept→company / **concept→concept**).
- 87.2% of edges have the filer ORG as source.
- 8,490,524 distinct (entity, relationship, target) triples → **~52% of rows are repeat observations**.

Queried live on FinReflectKG (this review): node/edge/chunk counts; Tim Cook's 5 edges and the named-exec company counts; POSITION = 9,984; PERSON role-verb counts; ~99% of PERSON edges carry startDate; the aapl → net income per-year counts; the cross-ticker 2-hop probe.

ArangoDB behaviour — official docs (3.12): disjoint SmartGraphs prohibit edges between different smartGraphAttribute values; smart nodes may connect to satellite nodes; a relation's two sides must each be all-smart or all-satellite.

ArangoDB behaviour — live spike on the 3.12.x cluster (this review):

- OneShard: DB sharding: "single", collections distributeShardsLike: "_graphs", single DBServer leader (the FinReflectKgOneShard build).
- Hybrid disjoint SmartGraph: the **smart↔satellite edge collection is smart-sharded** (distributeShardsLike the smart vertex collection, replicationFactor: 1); **only the satellite↔satellite edge collection is replicated** (replicationFactor: "satellite"). So Option B replicates the concept vertices (17%) + concept→concept edges (27.7%), not the company↔concept edges.
- VCI selectivity: the typed direct edge query for net income disclosers scans 59,315 edges (the FinReflectKgOneShard validate.py run).

Judgments / recommendations (not facts): that Option C is the best fit for the current use case; that the cross-company analytical suite belongs on OneShard/baseline; the reified role-node refinement. These are reasoned positions, open to challenge.